

In [6]: `import numpy as np`

```
a = np.eye(4)
print(a)
print("\n")

b = np.arange(5, 20, 3)
print(b)
print("\n")

c = np.linspace(10, 20, 6)
print(c)
print("\n")
+
d = np.random.rand(4)
print(d)
print("\n")

e = np.random.randint(10, 50, 6)
print(e)
```

```
[[1.  0.  0.  0.]
 [0.  1.  0.  0.]
 [0.  0.  1.  0.]
 [0.  0.  0.  1.]]
```

```
[ 5  8 11 14 17]
```

```
[10. 12. 14. 16. 18. 20.]
```

```
[0.46329099 0.49236886 0.72397955 0.38182704]
```

```
[45 42 19 43 28 38]
```

In [35]: *# array inspection functions*

```
import numpy as np

f = np.shape([[2, 4, 6], [8, 10, 12]])
print(f)

g = np.array([[9, 3, 7], [1, 5, 11]])
h = np.shape(g)
print(h)

i = np.size(g)
print(i)

j = np.ndim(g)
print(j)

print(type(g))
```

```
(2, 3)
(2, 3)
6
2
<class 'numpy.ndarray'>
```

```
In [3]: #array manipulation
k=np.transpose(g)
print(k)
```

```
[[ 9  1]
 [ 3  5]
 [ 7 11]]
```

```
In [4]: import numpy as np
```

```
k = np.array([[12, 4],
               [6, 10]])
l = np.array([[3, 2],
               [1, 5]])
m = np.add(k, l)
print(m)
n = np.subtract(k, l)
print(n)
o = np.multiply(k, l)
print(o)
p = np.divide(k, l)
print(p)
q = np.sqrt(l)
print(q)
r = np.log(l)
print(r)
```

```
[[15  6]
 [ 7 15]]
[[9 2]
 [5 5]]
[[36  8]
 [ 6 50]]
[[4. 2.]
 [6. 2.]]
[[1.73205081 1.41421356]
 [1.         2.23606798]]
[[1.09861229 0.69314718]
 [0.         1.60943791]]
```

```
In [7]: import numpy as np
```

```
# statistical functions
l = np.array([[12, 6],
               [3, 15]])

s = np.min(l)
print(s, "\n")

t = np.max(l)
print(t, "\n")

u = np.mean(l)
v = np.median(l)
w = np.std(l)
```

```
x = np.var(1)

print(u, "\n")
print(v, "\n")
print(w, "\n")
print(x, "\n")

y = np.array([10, 12, 14, 16, 18])
z = np.mean(y)
print(z)
```

3

15

9.0

9.0

4.743416490252569

22.5

14.0

```
In [8]: #sorting and searching funcs
a=np.array([95,48,90,5,4,3,66,2])
np.sort(a)
```

```
Out[8]: array([ 2,  3,  4,  5, 48, 66, 90, 95])
```

```
In [9]: ab=np.array([8,8,9,4,6,6,1,1])
np.unique(ab)
```

```
Out[9]: array([1, 4, 6, 8, 9])
```

```
In [13]: import pandas as pd
bb=pd.read_csv("iris.csv")
bb.head()
```

```
Out[13]:
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
In [14]: bb.tail()
```

Out[14]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2
148	6.2	3.4	5.4	2.3	2
149	5.9	3.0	5.1	1.8	2

In [15]: `bb.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sepal length (cm)     150 non-null   float64
1   sepal width (cm)      150 non-null   float64
2   petal length (cm)     150 non-null   float64
3   petal width (cm)      150 non-null   float64
4   species               150 non-null   int64
dtypes: float64(4), int64(1)
memory usage: 6.0 KB
```

In [16]: `bb.describe()`

Out[16]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333	1.000000
std	0.828066	0.435866	1.765298	0.762238	0.819232
min	4.300000	2.000000	1.000000	0.100000	0.000000
25%	5.100000	2.800000	1.600000	0.300000	0.000000
50%	5.800000	3.000000	4.350000	1.300000	1.000000
75%	6.400000	3.300000	5.100000	1.800000	2.000000
max	7.900000	4.400000	6.900000	2.500000	2.000000

In [17]: `bb.shape`

Out[17]: (150, 5)

In [19]: `bb.columns`

```
Out[19]: Index(['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)',
               'petal width (cm)', 'species'],
              dtype='object')
```

In [22]: `df[['sepal length (cm)', 'sepal width (cm)']]`

Out[22]:

	sepal length (cm)	sepal width (cm)
0	5.1	3.5
1	4.9	3.0
2	4.7	3.2
3	4.6	3.1
4	5.0	3.6
...	...	...
145	6.7	3.0
146	6.3	2.5
147	6.5	3.0
148	6.2	3.4
149	5.9	3.0

150 rows × 2 columns

In [24]: `bb.loc[1, 'sepal length (cm)']`

Out[24]: `np.float64(4.9)`

In [25]: `#bb.drop()
#bb.rename()
#bb.replace()`

In [26]: `bb.to_csv("iris.csv")`

In [31]: `bb = df
bb.to_excel("bro.xlsx", index=False)`

In [32]: `bb.groupby('sepal length (cm)')`

Out[32]: `<pandas.core.groupby.generic.DataFrameGroupBy object at 0x00000281CB24BF80>`

In [33]: `import pandas as pd

c = pd.DataFrame({
 'Data': [10, 12, 14, 16, 18],
 'age': [22, 23, 25, 26, 28]
})

pd.concat([bb, c])`

Out[33]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species	Data	age
0	5.1	3.5	1.4	0.2	0.0	NaN	NaN
1	4.9	3.0	1.4	0.2	0.0	NaN	NaN
2	4.7	3.2	1.3	0.2	0.0	NaN	NaN
3	4.6	3.1	1.5	0.2	0.0	NaN	NaN
4	5.0	3.6	1.4	0.2	0.0	NaN	NaN
...	...	...	...	...	...	...	...
0	NaN	NaN	NaN	NaN	NaN	10.0	22.0
1	NaN	NaN	NaN	NaN	NaN	12.0	23.0
2	NaN	NaN	NaN	NaN	NaN	14.0	25.0
3	NaN	NaN	NaN	NaN	NaN	16.0	26.0
4	NaN	NaN	NaN	NaN	NaN	18.0	28.0

155 rows × 7 columns

In [36]: `bb.join(c)`

Out[36]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species	Data	age
0	5.1	3.5	1.4	0.2	0	10.0	22.0
1	4.9	3.0	1.4	0.2	0	12.0	23.0
2	4.7	3.2	1.3	0.2	0	14.0	25.0
3	4.6	3.1	1.5	0.2	0	16.0	26.0
4	5.0	3.6	1.4	0.2	0	18.0	28.0
...	...	...	...	...	...	...	...
145	6.7	3.0	5.2	2.3	2	NaN	NaN
146	6.3	2.5	5.0	1.9	2	NaN	NaN
147	6.5	3.0	5.2	2.0	2	NaN	NaN
148	6.2	3.4	5.4	2.3	2	NaN	NaN
149	5.9	3.0	5.1	1.8	2	NaN	NaN

150 rows × 7 columns

In [1]: *#Linear Regression*

```
In [14]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

# Your original basic dataset
X = np.array([1, 2, 3, 4, 5])
Y = np.array([3, 6, 7, 9, 10])

x_mean = np.mean(X)
y_mean = np.mean(Y)


num = np.sum((X - x_mean) * (Y - y_mean))
den = np.sum((X - x_mean) ** 2)
beta1 = num / den
beta0 = y_mean - beta1 * x_mean

print("Slope (beta1): ", beta1)
print("Intercept (beta0): ", beta0)

Y_pred = beta0 + beta1 * X

plt.scatter(X, Y, color="blue", label="Actual Data")
plt.plot(X, Y_pred, color="red", label="Regression Line")
plt.xlabel("X values")
plt.ylabel("Y values")
plt.title("Linear Regression on Basic Dataset")
plt.legend()
plt.show()
```

```
Slope (beta1):  1.7
Intercept (beta0):  1.9000000000000004
```



```
In [15]: # Your original Salary dataset (Kaggle: https://www.kaggle.com/datasets/abhishek)
# Save as Salary.csv
data = pd.read_csv('Salary_dataset(program1).csv')
X_salary = data["YearsExperience"].values
Y_salary = data["Salary"].values

x_mean_salary = np.mean(X_salary)
y_mean_salary = np.mean(Y_salary)


num_salary = np.sum((X_salary - x_mean_salary) * (Y_salary - y_mean_salary))
den_salary = np.sum((X_salary - x_mean_salary) ** 2)
beta1_salary = num_salary / den_salary
beta0_salary = y_mean_salary - beta1_salary * x_mean_salary

print("Slope (beta1):", beta1_salary)
print("Intercept (beta0):", beta0_salary)

Y_pred_salary = beta0_salary + beta1_salary * X_salary

plt.scatter(X_salary, Y_salary, color="blue", label="Actual Data")
plt.plot(X_salary, Y_pred_salary, color="red", label="Regression Line")
plt.xlabel("Years of Experience")
plt.ylabel("Salary")
plt.title("Linear Regression on Salary Dataset")
plt.legend()
plt.show()
```

Slope (beta1): 9449.962321455076
 Intercept (beta0): 24848.2039665232




```
In [16]: # Second dataset for Linear Regression (Kaggle: https://www.kaggle.com/datasets/)
# Download and save as housing.csv
housing = pd.read_csv('Housing_dataset(program1).csv')
X_housing = housing['area'].values      # feature
Y_housing = housing['price'].values     # target

# Same manual linear regression code as you used
x_mean_h = np.mean(X_housing)
y_mean_h = np.mean(Y_housing)
num_h = np.sum((X_housing - x_mean_h) * (Y_housing - y_mean_h))
den_h = np.sum((X_housing - x_mean_h) ** 2)
beta1_h = num_h / den_h
beta0_h = y_mean_h - beta1_h * x_mean_h

Y_pred_h = beta0_h + beta1_h * X_housing

plt.scatter(X_housing, Y_housing, color="blue", label="Actual Data")
plt.plot(X_housing, Y_pred_h, color="red", label="Regression Line")
plt.xlabel("Area")
plt.ylabel("Price")
plt.title("Linear Regression on Housing Prices Dataset")
plt.legend()
plt.show()
```



In [17]: *#Logistic Regression*

```
In [18]: X_log = np.array([1, 2, 3, 4, 5])
y_log = np.array([0, 0, 0, 1, 1])

w = 0.0
b = 0.0
lr = 0.1
epochs = 1000

def sigmoid(z):
    return 1 / (1 + np.exp(-z))

for _ in range(epochs):
    z = w * X_log + b
    y_pred = sigmoid(z)
    dw = np.mean((y_pred - y_log) * X_log)
    db = np.mean(y_pred - y_log)
    w = w - lr * dw
    b = b - lr * db

print("Weight:", w)
print("Bias:", b)

def predict_logistic(X):
    probs = sigmoid(w * X + b)
    return np.where(probs >= 0.5, 1, 0)

print("Predicted Output:", predict_logistic(X_log))
```

Weight: 1.769045864973106
Bias: -5.956350539118638
Predicted Output: [0 0 0 1 1]

```
In [19]: # Logistic Regression - Heart Disease (Kaggle: https://www.kaggle.com/datasets/j
heart = pd.read_csv('heart_dataset(program1).csv')
X_heart = heart.drop('target', axis=1).values
y_heart = heart['target'].values

# Simple Logistic Regression with sklearn (standard way)
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

X_train, X_test, y_train, y_test = train_test_split(X_heart, y_heart, test_size=
logreg = LogisticRegression(max_iter=1000)
logreg.fit(X_train, y_train)
y_pred_heart = logreg.predict(X_test)
print("Heart Logistic Accuracy:", accuracy_score(y_test, y_pred_heart))
```

Heart Logistic Accuracy: 0.8051948051948052

```
In [20]: # Logistic Regression - Pima Diabetes (Kaggle: https://www.kaggle.com/datasets/u
diabetes = pd.read_csv('diabetes_ds(program1).csv')
X_dia = diabetes.drop('Outcome', axis=1).values
y_dia = diabetes['Outcome'].values

X_train, X_test, y_train, y_test = train_test_split(X_dia, y_dia, test_size=0.3,
logreg = LogisticRegression(max_iter=1000)
logreg.fit(X_train, y_train)
y_pred_dia = logreg.predict(X_test)
print("Diabetes Logistic Accuracy:", accuracy_score(y_test, y_pred_dia))
```

Diabetes Logistic Accuracy: 0.7359307359307359

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_rep

import warnings
warnings.filterwarnings('ignore')
sns.set_style("whitegrid")
```

```
In [2]: # Iris dataset
from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['target'] = iris.target

print(df.head())
print("\nDataset shape:", df.shape)
print("\nDataset Info:")
print(df.info())
print("\nStatistical Summary:")
print(df.describe())

# Check missing values
print("\nMissing values:")
print(df.isnull().sum())

sns.pairplot(df, hue="target")
plt.show()

# Correlation heatmap
plt.figure(figsize=(8,6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Matrix - Iris")
plt.show()

# Prepare data
X = df.drop('target', axis=1)
y = df['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_

# Scale (for KNN & SVM)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# 1. Decision Tree
dt = DecisionTreeClassifier(criterion='entropy', random_state=42)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
print("\nDecision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
```

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_dt))
print(classification_report(y_test, y_pred_dt))

# 2. KNN
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_scaled, y_train)
y_pred_knn = knn.predict(X_test_scaled)
print("\nKNN Accuracy:", accuracy_score(y_test, y_pred_knn))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_knn))
print(classification_report(y_test, y_pred_knn))

# 3. Naive Bayes
nb = GaussianNB()
nb.fit(X_train, y_train)
y_pred_nb = nb.predict(X_test)
print("\nNaive Bayes Accuracy:", accuracy_score(y_test, y_pred_nb))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_nb))
print(classification_report(y_test, y_pred_nb))

# 4. SVM
svm = SVC(kernel='linear', random_state=42)
svm.fit(X_train_scaled, y_train)
y_pred_svm = svm.predict(X_test_scaled)
print("\nSVM Accuracy:", accuracy_score(y_test, y_pred_svm))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_svm))
print(classification_report(y_test, y_pred_svm))
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	target
0	0
1	0
2	0
3	0
4	0

Dataset shape: (150, 5)

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 150 entries, 0 to 149

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	sepal length (cm)	150 non-null	float64
1	sepal width (cm)	150 non-null	float64
2	petal length (cm)	150 non-null	float64
3	petal width (cm)	150 non-null	float64
4	target	150 non-null	int64

dtypes: float64(4), int64(1)

memory usage: 6.0 KB

None

Statistical Summary:


	sepal length (cm)	sepal width (cm)	petal length (cm)	\
count	150.000000	150.000000	150.000000	
mean	5.843333	3.057333	3.758000	
std	0.828066	0.435866	1.765298	
min	4.300000	2.000000	1.000000	
25%	5.100000	2.800000	1.600000	
50%	5.800000	3.000000	4.350000	
75%	6.400000	3.300000	5.100000	
max	7.900000	4.400000	6.900000	


	petal width (cm)	target
count	150.000000	150.000000
mean	1.199333	1.000000
std	0.762238	0.819232
min	0.100000	0.000000
25%	0.300000	0.000000
50%	1.300000	1.000000
75%	1.800000	2.000000
max	2.500000	2.000000

Missing values:

sepal length (cm)	0
sepal width (cm)	0
petal length (cm)	0
petal width (cm)	0
target	0

dtype: int64





Decision Tree Accuracy: 0.9777777777777777

Confusion Matrix:

```
[[19  0  0]
 [ 0 13  0]
 [ 0  1 12]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	0.93	1.00	0.96	13
2	1.00	0.92	0.96	13
accuracy			0.98	45
macro avg	0.98	0.97	0.97	45
weighted avg	0.98	0.98	0.98	45

KNN Accuracy: 1.0

Confusion Matrix:

```
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

Naive Bayes Accuracy: 0.9777777777777777

Confusion Matrix:

```
[[19  0  0]
 [ 0 12  1]
 [ 0  0 13]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	0.92	0.96	13
2	0.93	1.00	0.96	13
accuracy			0.98	45
macro avg	0.98	0.97	0.97	45
weighted avg	0.98	0.98	0.98	45

SVM Accuracy: 0.9777777777777777

Confusion Matrix:

```
[[19  0  0]
 [ 0 12  1]
 [ 0  0 13]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	0.92	0.96	13
2	0.93	1.00	0.96	13
accuracy			0.98	45

macro avg	0.98	0.97	0.97	45
weighted avg	0.98	0.98	0.98	45

```
In [6]: # Heart Disease dataset
# Load dataset (Kaggle: johnsmith88/heart-disease-dataset)
heart = pd.read_csv('heart_ds.csv')

print(heart.head())
print("\nDataset shape:", heart.shape)
print("\nDataset Info:")
print(heart.info())
print("\nStatistical Summary:")
print(heart.describe())

# Missing values
print("\nMissing values:")
print(heart.isnull().sum())

# Pairplot
sns.pairplot(heart, hue="target")
plt.show()

# Correlation heatmap
plt.figure(figsize=(10,8))
sns.heatmap(heart.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Matrix - Heart Disease")
plt.show()

# Prepare data
X = heart.drop('target', axis=1)
y = heart['target']

# Label encoding (in case any object columns - matches your program4 style)
for col in X.select_dtypes(include=['object']).columns:
    le = LabelEncoder()
    X[col] = le.fit_transform(X[col])

# Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_

# Scale
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Decision Tree
dt = DecisionTreeClassifier(criterion='entropy', random_state=42)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
print("\nDecision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_dt))
print(classification_report(y_test, y_pred_dt))

# KNN
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_scaled, y_train)
y_pred_knn = knn.predict(X_test_scaled)
print("\nKNN Accuracy:", accuracy_score(y_test, y_pred_knn))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_knn))
```

```
print(classification_report(y_test, y_pred_knn))

# Naive Bayes
nb = GaussianNB()
nb.fit(X_train, y_train)
y_pred_nb = nb.predict(X_test)
print("\nNaive Bayes Accuracy:", accuracy_score(y_test, y_pred_nb))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_nb))
print(classification_report(y_test, y_pred_nb))

# SVM
svm = SVC(kernel='linear', random_state=42)
svm.fit(X_train_scaled, y_train)
y_pred_svm = svm.predict(X_test_scaled)
print("\nSVM Accuracy:", accuracy_score(y_test, y_pred_svm))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_svm))
print(classification_report(y_test, y_pred_svm))
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	
4	62	0	0	138	294	1	1	106	0	1.9	1	

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

Dataset shape: (1025, 14)

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1025 entries, 0 to 1024

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	age	1025 non-null	int64
1	sex	1025 non-null	int64
2	cp	1025 non-null	int64
3	trestbps	1025 non-null	int64
4	chol	1025 non-null	int64
5	fbs	1025 non-null	int64
6	restecg	1025 non-null	int64
7	thalach	1025 non-null	int64
8	exang	1025 non-null	int64
9	oldpeak	1025 non-null	float64
10	slope	1025 non-null	int64
11	ca	1025 non-null	int64
12	thal	1025 non-null	int64
13	target	1025 non-null	int64

dtypes: float64(1), int64(13)

memory usage: 112.2 KB

None

Statistical Summary:

	age	sex	cp	trestbps	chol	\
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	
mean	54.434146	0.695610	0.942439	131.611707	246.000000	
std	9.072290	0.460373	1.029641	17.516718	51.59251	
min	29.000000	0.000000	0.000000	94.000000	126.000000	
25%	48.000000	0.000000	0.000000	120.000000	211.000000	
50%	56.000000	1.000000	1.000000	130.000000	240.000000	
75%	61.000000	1.000000	2.000000	140.000000	275.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	


	fbs	restecg	thalach	exang	oldpeak	\
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	
mean	0.149268	0.529756	149.114146	0.336585	1.071512	
std	0.356527	0.527878	23.005724	0.472772	1.175053	
min	0.000000	0.000000	71.000000	0.000000	0.000000	
25%	0.000000	0.000000	132.000000	0.000000	0.000000	
50%	0.000000	1.000000	152.000000	0.000000	0.800000	
75%	0.000000	1.000000	166.000000	1.000000	1.800000	
max	1.000000	2.000000	202.000000	1.000000	6.200000	


	slope	ca	thal	target
count	1025.000000	1025.000000	1025.000000	1025.000000
mean	1.385366	0.754146	2.323902	0.513171
std	0.617755	1.030798	0.620660	0.500070
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	2.000000	0.000000
50%	1.000000	0.000000	2.000000	1.000000
75%	2.000000	1.000000	3.000000	1.000000
max	2.000000	4.000000	3.000000	1.000000

Missing values:

age	0
sex	0
cp	0
trestbps	0
chol	0
fbs	0
restecg	0
thalach	0
exang	0
oldpeak	0
slope	0
ca	0
thal	0
target	0

dtype: int64





Decision Tree Accuracy: 0.9902597402597403

Confusion Matrix:

```
[[159  0]
 [  3 146]]
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	159
1	1.00	0.98	0.99	149
accuracy			0.99	308
macro avg	0.99	0.99	0.99	308
weighted avg	0.99	0.99	0.99	308

KNN Accuracy: 0.8409090909090909

Confusion Matrix:

```
[[130  29]
 [ 20 129]]
```

	precision	recall	f1-score	support
0	0.87	0.82	0.84	159
1	0.82	0.87	0.84	149
accuracy			0.84	308
macro avg	0.84	0.84	0.84	308
weighted avg	0.84	0.84	0.84	308

Naive Bayes Accuracy: 0.814935064935065

Confusion Matrix:

```
[[118  41]
 [ 16 133]]
```

	precision	recall	f1-score	support
0	0.88	0.74	0.81	159
1	0.76	0.89	0.82	149
accuracy			0.81	308
macro avg	0.82	0.82	0.81	308
weighted avg	0.82	0.81	0.81	308

SVM Accuracy: 0.8084415584415584

Confusion Matrix:

```
[[119  40]
 [ 19 130]]
```

	precision	recall	f1-score	support
0	0.86	0.75	0.80	159
1	0.76	0.87	0.82	149
accuracy			0.81	308
macro avg	0.81	0.81	0.81	308
weighted avg	0.82	0.81	0.81	308

```
In [4]: #PIMA INDIANS DIABETES DATASET
# Load dataset (Kaggle: uciml/pima-indians-diabetes-database)
diabetes = pd.read_csv('diabetes_ds.csv')
```



```
print(diabetes.head())
print("\nDataset shape:", diabetes.shape)
print("\nDataset Info:")
print(diabetes.info())
print("\nStatistical Summary:")
print(diabetes.describe())

# Missing values
print("\nMissing values:")
print(diabetes.isnull().sum())

# Pairplot
sns.pairplot(diabetes, hue="Outcome")
plt.show()

# Correlation heatmap
plt.figure(figsize=(10,8))
sns.heatmap(diabetes.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Matrix - Diabetes")
plt.show()

# Prepare data
X = diabetes.drop('Outcome', axis=1)
y = diabetes['Outcome']

# Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_

# Scale
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Decision Tree
dt = DecisionTreeClassifier(criterion='entropy', random_state=42)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
print("\nDecision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_dt))
print(classification_report(y_test, y_pred_dt))

# KNN
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_scaled, y_train)
y_pred_knn = knn.predict(X_test_scaled)
print("\nKNN Accuracy:", accuracy_score(y_test, y_pred_knn))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_knn))
print(classification_report(y_test, y_pred_knn))

# Naive Bayes
nb = GaussianNB()
nb.fit(X_train, y_train)
y_pred_nb = nb.predict(X_test)
print("\nNaive Bayes Accuracy:", accuracy_score(y_test, y_pred_nb))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_nb))
print(classification_report(y_test, y_pred_nb))

# SVM
svm = SVC(kernel='linear', random_state=42)
svm.fit(X_train_scaled, y_train)
```

```
y_pred_svm = svm.predict(X_test_scaled)
print("\nSVM Accuracy:", accuracy_score(y_test, y_pred_svm))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_svm))
print(classification_report(y_test, y_pred_svm))
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Dataset shape: (768, 9)

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 768 entries, 0 to 767

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

dtypes: float64(2), int64(7)

memory usage: 54.1 KB

None

Statistical Summary:


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000


	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

Missing values:

Pregnancies	0
Glucose	0
BloodPressure	0

SkinThickness 0
Insulin 0
BMI 0
DiabetesPedigreeFunction 0
Age 0
Outcome 0
dtype: int64





Decision Tree Accuracy: 0.7272727272727273

Confusion Matrix:

```
[[118 33]
```

```
[ 30 50]]
```

	precision	recall	f1-score	support
0	0.80	0.78	0.79	151
1	0.60	0.62	0.61	80
accuracy			0.73	231
macro avg	0.70	0.70	0.70	231
weighted avg	0.73	0.73	0.73	231

KNN Accuracy: 0.7012987012987013

Confusion Matrix:

```
[[119 32]
```

```
[ 37 43]]
```

	precision	recall	f1-score	support
0	0.76	0.79	0.78	151
1	0.57	0.54	0.55	80
accuracy			0.70	231
macro avg	0.67	0.66	0.67	231
weighted avg	0.70	0.70	0.70	231

Naive Bayes Accuracy: 0.7445887445887446

Confusion Matrix:

```
[[119 32]
```

```
[ 27 53]]
```

	precision	recall	f1-score	support
0	0.82	0.79	0.80	151
1	0.62	0.66	0.64	80
accuracy			0.74	231
macro avg	0.72	0.73	0.72	231
weighted avg	0.75	0.74	0.75	231

SVM Accuracy: 0.7489177489177489

Confusion Matrix:

```
[[123 28]
```

```
[ 30 50]]
```

	precision	recall	f1-score	support
0	0.80	0.81	0.81	151
1	0.64	0.62	0.63	80
accuracy			0.75	231
macro avg	0.72	0.72	0.72	231
weighted avg	0.75	0.75	0.75	231

In []:

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans, AgglomerativeClustering
from scipy.cluster.hierarchy import dendrogram, linkage

sns.set_style("whitegrid")
```

```
In [ ]: #K-Means & Hierarchical Clustering

#MALL CUSTOMERS DATASET
df = pd.read_csv('Mall_Customers.csv') # Kaggle Link https://www.kaggle.com/data
print(df.head(), "\n\n")
print(df.info())

print("\nMissing Values:")
print(df.isnull().sum())

df.dropna(inplace=True)
data = df[['Annual Income (k$)', 'Spending Score (1-100)']]
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40


```
<class 'pandas.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                200 non-null   str
2   Age                  200 non-null   int64
3   Annual Income (k$)    200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), str(1)
memory usage: 7.9 KB
None
```

```
Missing Values:
CustomerID      0
Gender          0
Age             0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64
```

```
In [ ]: plt.figure(figsize=(8,6))
plt.scatter(data['Annual Income (k$)'], data['Spending Score (1-100)'])
plt.xlabel("Annual Income")
plt.ylabel("Spending Score")
plt.title("Customer Data Distribution")
plt.show()
```

```
wcss = []  
for i in range(1,11):  
    kmeans = KMeans(n_clusters=i, random_state=42)  
    kmeans.fit(data)  
    wcss.append(kmeans.inertia_)  
  
plt.plot(range(1,11), wcss)  
plt.title("Elbow Method")  
plt.xlabel("Number of clusters")  
plt.ylabel("WCSS")  
plt.show()
```







```
In [5]: # K-Means Clustering
kmeans = KMeans(n_clusters=5, random_state=42)
clusters = kmeans.fit_predict(data)
data["Cluster"] = clusters


plt.figure(figsize=(7,6))
sns.scatterplot(x='Annual Income (k$)', y='Spending Score (1-100)', hue="Cluster")
plt.title("K-Means Clustering")
plt.show()

# Hierarchical Clustering - Dendrogram (exact from PDF)
linked = linkage(data[['Annual Income (k$)', 'Spending Score (1-100)']], method='ward', metric='euclidean')
plt.figure(figsize=(8,6))
dendrogram(linked)
plt.title("Dendrogram")
plt.show()

# Agglomerative Hierarchical Clustering (exact from PDF)
hc = AgglomerativeClustering(n_clusters=5)
data['hc_cluster'] = hc.fit_predict(data[['Annual Income (k$)', 'Spending Score (1-100)']])

plt.figure(figsize=(7,6))
sns.scatterplot(x='Annual Income (k$)', y='Spending Score (1-100)', hue='hc_cluster')
plt.title("Hierarchy Clustering")
plt.show()
```





```
In [7]: #WHOLESALE CUSTOMERS DATASET
wholesale = pd.read_csv('Wholesale_customers_data.csv')
print(wholesale.head(), "\n\n")
print(wholesale.info())

print("\nMissing Values:")
print(wholesale.isnull().sum())

# Using two features (mirroring Mall_Customers)
data2 = wholesale[['Fresh', 'Milk']]
print("\nSelected features for clustering: Fresh & Milk")
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844
3	1	3	13265	1196	4221	6404	507	1788
4	2	3	22615	5410	7198	3915	1777	5185

```
<class 'pandas.DataFrame'>
```

```
RangeIndex: 440 entries, 0 to 439
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	Channel	440 non-null	int64
1	Region	440 non-null	int64
2	Fresh	440 non-null	int64
3	Milk	440 non-null	int64
4	Grocery	440 non-null	int64
5	Frozen	440 non-null	int64
6	Detergents_Paper	440 non-null	int64
7	Delicassen	440 non-null	int64

```
dtypes: int64(8)
```

```
memory usage: 27.6 KB
```

```
None
```

```
Missing Values:
```

Channel	0
Region	0
Fresh	0
Milk	0
Grocery	0
Frozen	0
Detergents_Paper	0
Delicassen	0



```
dtype: int64
```

```
Selected features for clustering: Fresh & Milk
```

```
In [8]: # Scatter plot
plt.figure(figsize=(8,6))
plt.scatter(data2['Fresh'], data2['Milk'])
plt.xlabel("Fresh")
plt.ylabel("Milk")
plt.title("Wholesale Customers Data Distribution")
plt.show()

# Elbow Method
wcss2 = []
for i in range(1,11):
    kmeans = KMeans(n_clusters=i, random_state=42)
    kmeans.fit(data2)
    wcss2.append(kmeans.inertia_)

plt.plot(range(1,11), wcss2)
plt.title("Elbow Method - Wholesale Customers")
plt.xlabel("Number of clusters")
plt.ylabel("WCSS")
plt.show()
```




```
In [10]: # K-Means Clustering
kmeans2 = KMeans(n_clusters=5, random_state=42)
clusters2 = kmeans2.fit_predict(data2)
data2["Cluster"] = clusters2
```



```
plt.figure(figsize=(7,6))
sns.scatterplot(x='Fresh', y='Milk', hue="Cluster", data=data2, palette='tab10')
plt.title("K-Means Clustering - Wholesale Customers")
plt.show()

# Hierarchical - Dendrogram
linked2 = linkage(data2[['Fresh', 'Milk']], method='ward')
plt.figure(figsize=(8,6))
dendrogram(linked2)
plt.title("Dendrogram - Wholesale Customers")
plt.show()

# Agglomerative Hierarchical
hc2 = AgglomerativeClustering(n_clusters=5)
data2['hc_cluster'] = hc2.fit_predict(data2[['Fresh', 'Milk']])

plt.figure(figsize=(7,6))
sns.scatterplot(x='Fresh', y='Milk', hue='hc_cluster', data=data2, palette='tab10')
plt.title("Hierarchy Clustering - Wholesale Customers")
plt.show()
```





```
In [1]: #Ensemble Techniques (Bagging, Random Forest, Boosting)

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import BaggingClassifier, RandomForestClassifier, AdaBoost
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
```

```
In [3]: #Same diabetes dataset from Kaggle (https://www.kaggle.com/datasets/uciml/pima-i
data = pd.read_csv('diabetes_ds.csv')
print("Dataset Shape:", data.shape)
print(data.head())

X = data.drop("Outcome", axis=1)
y = data["Outcome"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
```

Dataset Shape: (768, 9)

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
In [4]: # Bagging Classifier
bagging = BaggingClassifier(
    estimator=DecisionTreeClassifier(),
    n_estimators=50,
    random_state=42
)
bagging.fit(X_train, y_train)
y_pred_bagg = bagging.predict(X_test)
print("Bagging accuracy :", accuracy_score(y_test, y_pred_bagg))

# Random Forest (exact from your program6.pdf)
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print("Random Forest accuracy :", accuracy_score(y_test, y_pred_rf))

# AdaBoost (exact from your program6.pdf)
boost = AdaBoostClassifier(n_estimators=50, random_state=42)
boost.fit(X_train, y_train)
y_pred_boos = boost.predict(X_test)
print("Boosting accuracy :", accuracy_score(y_test, y_pred_boos))
```


Bagging accuracy : 0.7467532467532467
Random Forest accuracy : 0.7207792207792207
Boosting accuracy : 0.7792207792207793


```
In [5]: #comparison of accuracies
bag_acc = accuracy_score(y_test, y_pred_bagg)
rf_acc = accuracy_score(y_test, y_pred_rf)
boo_acc = accuracy_score(y_test, y_pred_boos)


models = ['Bagging', 'RandomForest', 'Boosting']
accuracies = [bag_acc, rf_acc, boo_acc]

plt.figure(figsize=(8,5))
plt.bar(models, accuracies, color=['blue', 'orange', 'green'])
plt.title("Accuracy Comparison of Ensemble Models")
plt.xlabel("Models")
plt.ylabel("Accuracy")
plt.ylim(0.7, 0.85)
plt.show()

n_values = [10, 50, 100, 200]
rf_scores = []
for n in n_values:
    model = RandomForestClassifier(n_estimators=n, random_state=42)
    model.fit(X_train, y_train)
    pred = model.predict(X_test)
    rf_scores.append(accuracy_score(y_test, pred))


plt.figure(figsize=(8,5))
plt.plot(n_values, rf_scores, marker='o', color='blue')
plt.title("Effect of n_estimators on Accuracy (Random Forest)")
plt.xlabel("Number of Trees")
plt.ylabel("Accuracy")
plt.grid(True)
plt.show()
```





```
In [6]: # Feature importance
importances = rf.feature_importances_
features = X.columns
fea_imp = pd.Series(importances, index=features)

plt.figure(figsize=(10,6))
fea_imp.sort_values().plot(kind='barh', color='skyblue')
plt.title("Feature Importance (Random Forest)")
plt.xlabel("Importance")
plt.ylabel("Features")
plt.show()
```



```
In [7]: # LEARNER SVM & Naïve Bayes as base

# Bagging
bagging_svm = BaggingClassifier(
```

```

    estimator=SVC(),
    n_estimators=50,
    random_state=42
)
bagging_svm.fit(X_train, y_train)
y_pred_bagg_svm = bagging_svm.predict(X_test)
print("Bagging (SVM base) Accuracy:", accuracy_score(y_test, y_pred_bagg_svm))

# Bagging with Naive Bayes as base Learner
bagging_nb = BaggingClassifier(
    estimator=GaussianNB(),
    n_estimators=50,
    random_state=42
)
bagging_nb.fit(X_train, y_train)
y_pred_bagg_nb = bagging_nb.predict(X_test)
print("Bagging (Naive Bayes base) Accuracy:", accuracy_score(y_test, y_pred_bagg_nb))

```

Bagging (SVM base) Accuracy: 0.7662337662337663

Bagging (Naive Bayes base) Accuracy: 0.7792207792207793

In [9]:

```

# Boosting with SVM as base Learner
boost_svm = AdaBoostClassifier(
    estimator=SVC(probability=True), # AdaBoost needs probability=True for SVM
    n_estimators=50,
    random_state=42,
)
boost_svm.fit(X_train, y_train)
y_pred_boost_svm = boost_svm.predict(X_test)
print("Boosting (SVM base) Accuracy:", accuracy_score(y_test, y_pred_boost_svm))

# Boosting with Naive Bayes as base Learner
boost_nb = AdaBoostClassifier(
    estimator=GaussianNB(),
    n_estimators=50,
    random_state=42
)
boost_nb.fit(X_train, y_train)
y_pred_boost_nb = boost_nb.predict(X_test)
print("Boosting (Naive Bayes base) Accuracy:", accuracy_score(y_test, y_pred_boost_nb))

```

Boosting (SVM base) Accuracy: 0.6428571428571429

Boosting (Naive Bayes base) Accuracy: 0.7467532467532467

In [11]:

```

# Final comparison table for all 7 models
comparison = {
    'Model': [
        'Bagging (DT base)', 'Random Forest', 'AdaBoost (DT base)',
        'Bagging (SVM base)', 'Bagging (NB base)',
        'AdaBoost (SVM base)', 'AdaBoost (NB base)'
    ],
    'Accuracy': [
        bag_acc, rf_acc, boo_acc,
        accuracy_score(y_test, y_pred_bagg_svm),
        accuracy_score(y_test, y_pred_bagg_nb),
        accuracy_score(y_test, y_pred_boost_svm),
        accuracy_score(y_test, y_pred_boost_nb)
    ]
}

comp_df = pd.DataFrame(comparison)


```

```
print("\n=== FINAL COMPARATIVE ANALYSIS ===")
print(comp_df.round(4))

# Optional: Bar chart for all 7 models
plt.figure(figsize=(12,6))
plt.bar(comp_df['Model'], comp_df['Accuracy'], color='teal')
plt.xticks(rotation=45, ha='right')
plt.title("Comparative Analysis - All Ensemble Variants")
plt.ylabel("Accuracy")
plt.ylim(0.6, 0.85)
plt.show()
```

=== FINAL COMPARATIVE ANALYSIS ===

	Model	Accuracy
0	Bagging (DT base)	0.7468
1	Random Forest	0.7208
2	AdaBoost (DT base)	0.7792
3	Bagging (SVM base)	0.7662
4	Bagging (NB base)	0.7792
5	AdaBoost (SVM base)	0.6429
6	AdaBoost (NB base)	0.7468



```
In [ ]: # PCA and LDA Using Digit Recognizer Dataset
```

```
In [1]: import pandas as pd
df = pd.read_csv('train(DigitRecognizer).csv')
print(df.head())
print(df.shape)

print(df.columns)

#checking for null values
print(df.isnull().sum())

x = df.drop('label', axis=1)
y = df['label']

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)

from sklearn.decomposition import PCA
# Here we choose to reduce to 2 dimensions for visualization purposes
pca_2 = PCA(n_components=2)
x_pca_2 = pca_2.fit_transform(x_scaled)
print("Explained variance ratio:", pca_2.explained_variance_ratio_)

import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
plt.scatter(x_pca_2[:, 0], x_pca_2[:, 1], c=y, cmap='tab10')
plt.colorbar()
plt.title('PCA on MNIST Dataset (Digit Recognizer) for 2 Components')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
lda_2 = LDA(n_components=2)
x_lda_2 = lda_2.fit_transform(x_scaled, y)

plt.figure(figsize=(8, 6))
plt.scatter(x_lda_2[:, 0], x_lda_2[:, 1], c=y, cmap='tab10')
plt.colorbar()
plt.title('LDA on MNIST Dataset (Digit Recognizer) for 2 Components')
plt.xlabel('Linear Discriminant 1')
plt.ylabel('Linear Discriminant 2')
plt.show()

from sklearn.decomposition import PCA
# Here we choose to reduce to 100 dimensions for visualization purposes
pca_100 = PCA(n_components=100)
x_pca_100 = pca_100.fit_transform(x_scaled)
print("Explained variance ratio:", pca_100.explained_variance_ratio_)

plt.figure(figsize=(8, 6))
plt.scatter(x_pca_100[:, 0], x_pca_100[:, 1], c=y, cmap='tab10')
plt.colorbar()
plt.title('PCA on MNIST Dataset (Digit Recognizer) for 100 Components')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
```

```
plt.show()

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
lda_5 = LDA(n_components=5)
x_lda_5 = lda_5.fit_transform(x_scaled, y)

import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
plt.scatter(x_lda_5[:, 0], x_lda_5[:, 1], c=y, cmap='tab10')
plt.colorbar()
plt.title('LDA on MNIST Dataset (Digit Recognizer) for 100 Components')
plt.xlabel('Linear Discriminant 1')
plt.ylabel('Linear Discriminant 2')
plt.show()
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	\
0	1	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	0	
2	1	0	0	0	0	0	0	0	0	
3	4	0	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	0	

	pixel8	...	pixel774	pixel775	pixel776	pixel777	pixel778	pixel779	\
0	0	...	0	0	0	0	0	0	
1	0	...	0	0	0	0	0	0	
2	0	...	0	0	0	0	0	0	
3	0	...	0	0	0	0	0	0	
4	0	...	0	0	0	0	0	0	

	pixel780	pixel781	pixel782	pixel783
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 785 columns]

(42000, 785)

Index(['label', 'pixel0', 'pixel1', 'pixel2', 'pixel3', 'pixel4', 'pixel5',
'pixel6', 'pixel7', 'pixel8',
...
'pixel774', 'pixel775', 'pixel776', 'pixel777', 'pixel778', 'pixel779',
'pixel780', 'pixel781', 'pixel782', 'pixel783'],
dtype='object', length=785)

label	0
pixel0	0
pixel1	0
pixel2	0
pixel3	0


..

pixel779	0
pixel780	0
pixel781	0
pixel782	0
pixel783	0


Length: 785, dtype: int64

Explained variance ratio: [0.05747953 0.04111691]

PCA on MNIST Dataset (Digit Recognizer) for 2 Components





LDA on MNIST Dataset (Digit Recognizer) for 2 Components



Explained variance ratio: [0.05747953 0.04111691 0.03782867 0.02939862 0.02556439
0.02229844
0.01952552 0.01771605 0.0156266 0.01424956 0.01361009 0.01222541
0.01135736 0.0111309 0.01050311 0.01012326 0.00951303 0.00934514
0.00907259 0.00885327 0.00838907 0.00812027 0.00775161 0.00752312
0.0072769 0.00698756 0.00690455 0.00664924 0.00630675 0.00616301
0.00610922 0.00597037 0.00577284 0.00573682 0.00564822 0.00546322
0.0053943 0.00524361 0.00504843 0.0048853 0.00482244 0.00475901
0.00460008 0.00457629 0.00449584 0.00446722 0.00443837 0.00436837
0.00432713 0.00427027 0.00419269 0.0041218 0.00402461 0.00399434
0.00394891 0.00390805 0.00379899 0.00372454 0.00368413 0.00365723
0.00353278 0.00351088 0.00345414 0.00341394 0.00337784 0.00336477
0.0033171 0.00329725 0.00320016 0.00316776 0.00312695 0.00311861
0.00308212 0.00303276 0.00301509 0.0029714 0.00294849 0.00293629
0.00287848 0.00286906 0.00283171 0.00282885 0.00281701 0.00279215
0.00276842 0.00276316 0.00273923 0.0027181 0.00268141 0.00265443
0.00262072 0.00260164 0.00258196 0.00254877 0.00253504 0.00252401
0.00250074 0.00245803 0.0024436 0.00241387]

PCA on MNIST Dataset (Digit Recognizer) for 100 Components





In []: *# PCA and LDA with Fashion-MNIST Dataset*

```
In [24]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

# 1. Load Data
df = pd.read_csv('fashion-mnist_train.csv')
print(df.head())
print(df.shape)

# 2. Check Columns and Nulls
print(df.columns)
print(df.isnull().sum())

# 3. Split Features and Target
x = df.drop('label', axis=1)
y = df['label']

# 4. Standard Scaling
scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)

# 5. PCA for 2 Components
pca_2 = PCA(n_components=2)
x_pca_2 = pca_2.fit_transform(x_scaled)
```

```
plt.figure(figsize=(8, 6))
plt.scatter(x_pca_2[:, 0], x_pca_2[:, 1], c=y, cmap='tab10')
plt.colorbar()
plt.title('PCA on Fashion MNIST for 2 Components')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()

# 6. LDA for 2 Components
lda_2 = LDA(n_components=2)
x_lda_2 = lda_2.fit_transform(x_scaled, y)

plt.figure(figsize=(8, 6))
plt.scatter(x_lda_2[:, 0], x_lda_2[:, 1], c=y, cmap='tab10')
plt.colorbar()
plt.title('LDA on Fashion MNIST for 2 Components')
plt.xlabel('Linear Discriminant 1')
plt.ylabel('Linear Discriminant 2')
plt.show()

# 7. PCA for 100 Components
pca_100 = PCA(n_components=100)
x_pca_100 = pca_100.fit_transform(x_scaled)

plt.figure(figsize=(8, 6))
plt.scatter(x_pca_100[:, 0], x_pca_100[:, 1], c=y, cmap='tab10')
plt.colorbar()
plt.title('PCA on Fashion MNIST (100 Components View)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()

# 8. LDA for 5 Components
lda_5 = LDA(n_components=5)
x_lda_5 = lda_5.fit_transform(x_scaled, y)

plt.figure(figsize=(8, 6))
plt.scatter(x_lda_5[:, 0], x_lda_5[:, 1], c=y, cmap='tab10')
plt.colorbar()
plt.title('LDA on Fashion MNIST for 5 Components')
plt.xlabel('Linear Discriminant 1')
plt.ylabel('Linear Discriminant 2')
plt.show()
```

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	\
0	2	0	0	0	0	0	0	0	0	
1	9	0	0	0	0	0	0	0	0	
2	6	0	0	0	0	0	0	0	5	
3	0	0	0	0	1	2	0	0	0	
4	3	0	0	0	0	0	0	0	0	

	pixel9	...	pixel775	pixel776	pixel777	pixel778	pixel779	pixel780	\
0	0	...	0	0	0	0	0	0	
1	0	...	0	0	0	0	0	0	
2	0	...	0	0	0	30	43	0	
3	0	...	3	0	0	0	0	1	
4	0	...	0	0	0	0	0	0	

	pixel781	pixel782	pixel783	pixel784
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0



[5 rows x 785 columns]



(60000, 785)

Index(['label', 'pixel1', 'pixel2', 'pixel3', 'pixel4', 'pixel5', 'pixel6',
'pixel7', 'pixel8', 'pixel9',
...,
'pixel775', 'pixel776', 'pixel777', 'pixel778', 'pixel779', 'pixel780',
'pixel781', 'pixel782', 'pixel783', 'pixel784'],
dtype='str', length=785)

label	0
pixel1	0
pixel2	0
pixel3	0
pixel4	0
..	
pixel780	0
pixel781	0
pixel782	0
pixel783	0
pixel784	0

Length: 785, dtype: int64





In [22]: *# PCA and LDA using Dry Bean Dataset*

```
In [26]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.decomposition import PCA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

df = pd.read_csv('train(Human_Activity_Recognition).csv')
print(df.head())
print(df.shape)

print(df.columns)
print(df.isnull().sum())

# In this dataset, the target column is usually named 'Activity'
x = df.drop(['Activity', 'subject'], axis=1) # Dropping subject ID as it's not a
y_raw = df['Activity']

# Encode the string labels (Walking, Sitting, etc.) into numbers for plotting
le = LabelEncoder()
y = le.fit_transform(y_raw)

scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)

# PCA for 2 Components
pca_2 = PCA(n_components=2)
x_pca_2 = pca_2.fit_transform(x_scaled)

plt.figure(figsize=(8, 6))
plt.scatter(x_pca_2[:, 0], x_pca_2[:, 1], c=y, cmap='viridis')
plt.colorbar()
plt.title('PCA on HAR Dataset for 2 Components')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()

# LDA for 2 Components
lda_2 = LDA(n_components=2)
x_lda_2 = lda_2.fit_transform(x_scaled, y)

plt.figure(figsize=(8, 6))
plt.scatter(x_lda_2[:, 0], x_lda_2[:, 1], c=y, cmap='viridis')
plt.colorbar()
plt.title('LDA on HAR Dataset for 2 Components')
plt.xlabel('Linear Discriminant 1')
plt.ylabel('Linear Discriminant 2')
plt.show()

# PCA for 100 Components
pca_100 = PCA(n_components=100)
x_pca_100 = pca_100.fit_transform(x_scaled)

plt.figure(figsize=(8, 6))
plt.scatter(x_pca_100[:, 0], x_pca_100[:, 1], c=y, cmap='viridis')
plt.colorbar()
plt.title('PCA on HAR Dataset (100 Components View)')
plt.xlabel('Principal Component 1')
```

```
plt.ylabel('Principal Component 2')
plt.show()

# LDA for 5 Components
lda_5 = LDA(n_components=5)
x_lda_5 = lda_5.fit_transform(x_scaled, y)

plt.figure(figsize=(8, 6))
plt.scatter(x_lda_5[:, 0], x_lda_5[:, 1], c=y, cmap='viridis')
plt.colorbar()
plt.title('LDA on HAR Dataset for 5 Components')
plt.xlabel('Linear Discriminant 1')
plt.ylabel('Linear Discriminant 2')
plt.show()
```

	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X \
0	0.288585	-0.020294	-0.132905	-0.995279
1	0.278419	-0.016411	-0.123520	-0.998245
2	0.279653	-0.019467	-0.113462	-0.995380
3	0.279174	-0.026201	-0.123283	-0.996091
4	0.276629	-0.016570	-0.115362	-0.998139

	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAcc-mad()-X	tBodyAcc-mad()-Y \
0	-0.983111	-0.913526	-0.995112	-0.983185
1	-0.975300	-0.960322	-0.998807	-0.974914
2	-0.967187	-0.978944	-0.996520	-0.963668
3	-0.983403	-0.990675	-0.997099	-0.982750
4	-0.980817	-0.990482	-0.998321	-0.979672

	tBodyAcc-mad()-Z	tBodyAcc-max()-X	...	fBodyBodyGyroJerkMag-kurtosis() \
0	-0.923527	-0.934724	...	-0.710304
1	-0.957686	-0.943068	...	-0.861499
2	-0.977469	-0.938692	...	-0.760104
3	-0.989302	-0.938692	...	-0.482845
4	-0.990441	-0.942469	...	-0.699205

	angle(tBodyAccMean,gravity)	angle(tBodyAccJerkMean,gravityMean) \
0	-0.112754	0.030400
1	0.053477	-0.007435
2	-0.118559	0.177899
3	-0.036788	-0.012892
4	0.123320	0.122542

	angle(tBodyGyroMean,gravityMean)	angle(tBodyGyroJerkMean,gravityMean) \
0	-0.464761	-0.018446
1	-0.732626	0.703511
2	0.100699	0.808529
3	0.640011	-0.485366
4	0.693578	-0.615971

	angle(X,gravityMean)	angle(Y,gravityMean)	angle(Z,gravityMean)	subject \
0	-0.841247	0.179941	-0.058627	1
1	-0.844788	0.180289	-0.054317	1
2	-0.848933	0.180637	-0.049118	1
3	-0.848649	0.181935	-0.047663	1
4	-0.847865	0.185151	-0.043892	1


	Activity
0	STANDING
1	STANDING
2	STANDING
3	STANDING
4	STANDING



[5 rows x 563 columns]
(7352, 563)


```
Index(['tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y', 'tBodyAcc-mean()-Z',
      'tBodyAcc-std()-X', 'tBodyAcc-std()-Y', 'tBodyAcc-std()-Z',
      'tBodyAcc-mad()-X', 'tBodyAcc-mad()-Y', 'tBodyAcc-mad()-Z',
      'tBodyAcc-max()-X',
      ...,
      'fBodyBodyGyroJerkMag-kurtosis()', 'angle(tBodyAccMean,gravity)',
      'angle(tBodyAccJerkMean,gravityMean)',
      'angle(tBodyGyroMean,gravityMean)',
      'angle(tBodyGyroJerkMean,gravityMean)', 'angle(X,gravityMean)',
```



```
    'angle(Y,gravityMean)', 'angle(Z,gravityMean)', 'subject', 'Activity'],  
    dtype='str', length=563)  
tBodyAcc-mean()-X      0  
tBodyAcc-mean()-Y      0  
tBodyAcc-mean()-Z      0  
tBodyAcc-std()-X       0  
tBodyAcc-std()-Y       0  
..  
angle(X,gravityMean)    0  
angle(Y,gravityMean)    0  
angle(Z,gravityMean)    0  
subject                 0  
Activity                0  
Length: 563, dtype: int64
```







In []: